

# Naive bayes implementation over Amazon fine food reviews dataset

```
In [13]: #IMPORTING RELEVANT LIBRARIES

%matplotlib inline

import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
import itertools

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics

from sklearn.cross_validation import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.cross_validation import cross_val_score
from collections import Counter

from sklearn import cross_validation
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import MultinomialNB
from sklearn.model_selection import KFold
from sklearn.naive_bayes import BernoulliNB
from imblearn.over_sampling import SMOTE
from prettytable import PrettyTable
from sklearn.model_selection import cross_validate
from random import *
```

```
In [2]: #Connecting to the SQL table
con = sqlite3.connect('final.sqlite')

#Reading data from the database

Data = pd.read_sql_query("""
SELECT *
FROM Reviews """,con)
Data.shape
```

Out[2]: (364171, 12)

```
In [3]: # Drop index column
Data.drop(columns=['index'],inplace=True)
```

```
In [4]: Data["Time"]=pd.to_datetime(Data.Time)
Data.head(5)
```

Out[4]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score
0	150524	0006641040	ACITT7DI6IDDL	shari zychinski	0	0	positive 1970-00:00:00.9393
1	150506	0006641040	A2IW4PEEKO2R0U	Tracy	1	1	positive 1970-00:00:01.1947
2	150507	0006641040	A1S4A3IQ2MU7V4	sally sue "sally sue"	1	1	positive 1970-00:00:01.1914
3	150508	0006641040	AZGXZ2UUK6X	Catherine Hallberg "(Kate)"	1	1	positive 1970-00:00:01.0760
4	150509	0006641040	A3CMRKGE0P909G	Teresa	3	4	positive 1970-00:00:01.0183

```
In [5]: #Setting Time column as index of the dataframe
Data.set_index("Time",inplace=True)

#Sampling the above data
Sorted=Data.sort_index()
```

```
In [6]: Sorted.head()
```

Out[6]:


Time	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score
1970-01-01 00:00:00.939340800	150524	0006641040	ACITT7DI6IDDL	shari zychinski	0	0	positive
1970-01-01 00:00:00.940809600	150501	0006641040	AJ46FKXOVC7NR	Nicholas A Mesiano	2	2	positive
1970-01-01 00:00:00.944092800	451856	B00004CXX9	AIUWLEQ1ADEG5	Elizabeth Medina	0	0	positive
1970-01-01 00:00:00.944438400	374359	B00004CI84	A344SMIA5JECGM	Vincent P. Ross	1	2	positive
1970-01-01 00:00:00.946857600	451855	B00004CXX9	AJH6LUC1UT1ON	The Phantom of the Opera	0	0	positive

```
In [7]: #Sampling the above data

Sampled_data=Sorted.sample(n=50000,replace='False')
Sample_sort=Sampled_data.sort_index()
Sample_sort.shape
```

Out[7]: (50000, 10)

```
In [8]: Class=Sorted["Score"]
label=Sample_sort["Score"]
sns.countplot(x="Score",data=Sorted,palette="hls")
plt.show()
plt.savefig("count_plot")
```



<Figure size 432x288 with 0 Axes>

```
In [9]: #Dropping the Score column which are the actual class labels of the dataset
Sorted.drop(columns=['Score'],inplace=True)
Sorted.shape
```

Out[9]: (364171, 9)

```
In [10]: #Dropping the Score column from the sampled_set which are the actual class labels of the dataset
Sample_sort.drop(columns=['Score'],inplace=True)
Sample_sort.shape
```

Out[10]: (50000, 9)

## Observations

- Here after all the text-preprocessing and the data-cleaning only 364k datapoints remained.
- I have taken a sample size of 50k and also the whole data out of the total population for the purpose of analyzing and studying the behaviour of the data by applying the Naive bayes algorithm.
- First I took the "TIME" column and set as the index of the new sampled dataframe and then sorted accordingly in ascending order since the data has a temporal nature.
- By setting the "SCORE" column as a class label for classifying the reviews as a positive and negative.
- By observing the above bar plot it is clear that the above dataset is highly imbalanced and this may cause problems in the future analysis.

## Function for Splitting the dataset into Train,CV and Test sets (70:30)

```
In [11]: def data_split(x,y):
#Splitting the model into 70:30 split of Training and Cross_validate split
X_1, X_test, y_1, y_test = train_test_split(X, Y, test_size=0.3,shuffle=False,random_state=None)

# split the train data set into cross validation train and cross validation test
X_tr, X_cv, y_tr, y_cv = train_test_split(X_1, y_1, test_size=0.3,shuffle=False,random_state=None)

return X_tr,y_tr,X_cv,y_cv,X_test,y_test
```

## Preparing and printing the Train,cv and test sets

```
In [12]: X=Sorted
Y=Class

X_tr,y_tr,X_cv,y_cv,X_test,y_test=data_split(X,Y)

print("The shape of x_train is:",X_tr.shape)
print("the shape of y_train is:",y_tr.shape)
print("the shape of x_cv is:",X_cv.shape)
print("the shape of y_cv is:",y_cv.shape)
print("the shape of x_test is:",X_test.shape)
print("the shape of y_test is:",y_test.shape)
```

The shape of x\_train is: (178443, 9)  
the shape of y\_train is: (178443,)  
the shape of x\_cv is: (76476, 9)  
the shape of y\_cv is: (76476,)  
the shape of x\_test is: (109252, 9)  
the shape of y\_test is: (109252,)

## Utility function for training and calculating the missclassification error of the model

```
In [14]: #function for training the model
def train(X_tr,Y_tr,X_cv,y_cv):

    clf = MultinomialNB()
    clf.fit(X_tr, y_tr)

    pred= clf.predict(X_cv)

    acc = accuracy_score(y_cv, pred, normalize=True) * float(100)

    print('\n The train accuracy by using default alpha over cv set is = %f%% ' % ( acc))
    return pred,acc

def test(X_tr,y_tr,X_test,y_test):
    Nb = MultinomialNB()
    Nb.fit(X_tr,y_tr)
    pred = Nb.predict(X_test)
    acc = accuracy_score(y_test, pred, normalize=True)*float(100)
    print('\n**** The Test accuracy by using default alpha is %d%%' % (acc))

def optimal_test(optimal_a,X_tr,y_tr,X_cv,y_cv,X_test,y_test):
    # instantiate learning model a = optimal_alpha
    B_optimal =MultinomialNB(alpha=optimal_a)

    # fitting the model
    B_optimal.fit(X_tr, y_tr)

    # predict the response
    pred = B_optimal.predict(X_test)

    Y_pred=B_optimal.predict(X_cv)

    TRAIN_acc = accuracy_score(y_cv, Y_pred, normalize=True) * float(100)
    print("\n The train accuracy of the NB classifier for the best alpha=%f%% is %f%%"%(optimal_a,TRAIN_acc))

    print(""*100)

    # evaluate accuracy
    Acc = accuracy_score(y_test, pred,normalize=True) * float(100)
    print('\nThe Test accuracy of the NB classifier by using the best alpha = %f%% is %f%%' % (optimal_a, Acc))

    return pred, Acc

def cross_validation(X_cv,y_cv):
    a=[0.001, 0.0001, 0.12, 0.42,0.25,0.50,0.33,0.0065,0.72, 1]
    multinom= a
    cv_Scores=[]
    for a in multinom:
        clf =MultinomialNB(a)
        scores = cross_val_score(clf, X_cv, y_cv, cv=10,scoring='accuracy')
        cv_Scores.append(scores.mean())
    #printing the 10 Cross-Validation scores
    print(cv_Scores)
    return cv_Scores,multinom
```

## Utility function for visualizing the model scores and CV error of the model

```

In [15]: from sklearn.metrics import confusion_matrix
def Confusion_metric(y_test,y_pred,acc):
    print(metrics.confusion_matrix(y_test,y_pred))
    confusion=metrics.confusion_matrix(y_test,y_pred)

    plt.figure(figsize=(9,9))
    sns.heatmap(confusion, annot=True, fmt=".3f", linewidths=.5, square = True, cmap = 'Blues_r');
    plt.ylabel('Predicted label');
    plt.xlabel('Actual label');
    all_sample_title = 'Accuracy Score: {0}'.format( acc)
    plt.title(all_sample_title, size = 15);
    plt.show()

#Storing the values of the confusion matrix
    TN=confusion[0,0]
    FP=confusion[0,1]
    FN=confusion[1,0]
    TP=confusion[1,1]

# use float to perform true division, not integer division
    Class_acc=((TP + TN) / float(TP + TN + FP + FN))*100

#Code for classification error

    classification_error = ((FP + FN) / float(TP + TN + FP + FN))*100

#Code for finding the TPR,FPR,TNR,FNR

    TPR = (TP / float(FN + TP))*100

    FNR = (FN / float(FN + TP))*100

    TNR=(TN / float(TN + FP))*100

    FPR=(FP / float(TN + FP))*100

#Code for finding the Precision,Recall & F1_score

    precision = (TP/float(TP+FP))*100

    recall= (TP / float(FN + TP))*100

    F1_s= ((float(precision*recall)/float(precision+recall))*2)

    print()

    ptable=PrettyTable()
    ptable.title="The performance metrics of the above model are as follows: "
    ptable.field_names=["Metrics","Scores"]
    ptable.add_row(["Classification_accuracy",Class_acc])
    ptable.add_row(["Classification_error",classification_error])
    ptable.add_row(["True positive",TP])
    ptable.add_row(["False positive",FP])
    ptable.add_row(["True negative",TN])
    ptable.add_row(["False negative",FN])
    ptable.add_row(["True positive rate",TPR])
    ptable.add_row(["False negative rate",FNR])
    ptable.add_row(["True negative rate",TNR])
    ptable.add_row(["False positive rate",FPR])
    ptable.add_row(["Precision value",precision])
    ptable.add_row(["Recall value",recall])
    ptable.add_row(["f1_score value",F1_s])

    print(ptable)

#Function for plotting the missclassification errors

def MSE_plot(CV_scores,multinom):
#Calculating the minimum Missclassification error of the above model

    MSE = [1 - x for x in CV_scores]

# determining best a
    optimal_a =multinom[MSE.index(min(MSE))]
    print('\nThe optimal number of alpha value is %f%.' % optimal_a)

# plot misclassification error vs a
    plt.plot(multinom, MSE)

    for xy in zip(multinom, np.round(MSE,3)):
        plt.annotate('%s, %s' % xy, xy=xy, textcoords='data')

    plt.xlabel('Number of alpha')
    plt.ylabel('Misclassification Error')
    plt.show()

```

```
print("the misclassification error for alpha value is : ", np.round(MSE,3))
return optimal_a
```

## Utility function for vectorizing the given data & finding the top features

```
In [16]: #Function for vectorizing the train data

from sklearn.preprocessing import StandardScaler
def vec_train(vect,X_tr):
    import warnings
    warnings.filterwarnings("ignore")

    count_vect = vect #in scikit-learn
    BOW = count_vect.fit_transform(X_tr.values)

    return count_vect,BOW

#Function for vectorizing the CV data

def vec_cv(count,X_cv):
    cv=count.transform(X_cv.values)
    cv.get_shape()

    return cv

#Function for vectorizing the test data

def vec_test(count,X_test):
    test=count.transform(X_test.values)
    test.get_shape()

    return test

#Funtion for printing the total number of top features
def top_tfidf_feats(name,row, features, top_n=25):
    ''' Get top n tfidf values in row and return them with their corresponding feature names.'''
    topn_ids = np.argsort(row)[::-1][:top_n]
    top_feats = [(features[i], row[i]) for i in topn_ids]
    df = pd.DataFrame(top_feats)
    df.columns = ['feature', name]
    return df
```

## Bag of words vectorization technique

```
In [17]: #Initializing the count vectorizer
Count_vect=CountVectorizer(binary=True)

#vectorizing the X_train set
count,x_tr=vec_train(Count_vect,X_tr["CleanedText"])

print("The shape of the X_train is: ",x_tr.shape)

#Vectgorizing the X_crossvalidation set
x_cv=vec_cv(count,X_cv["CleanedText"])
print("The shape of the X_cv is: ",x_cv.shape)

#Vectorizing the X_test set
x_test=vec_test(count,X_test["CleanedText"])
print("The shape of the X_test is: ",x_test.shape)

#Printing the total length of the features
print("\nTop 25 feaures according to the Bow score are as follows")
features = Count_vect.get_feature_names()
len(features)

top_Bow = top_tfidf_feats("bow",x_tr[1,:].toarray()[0],features,25)
top_Bow
```

The shape of the X\_train is: (178443, 49468)  
The shape of the X\_cv is: (76476, 49468)  
The shape of the X\_test is: (109252, 49468)

Top 25 feaures according to the Bow score are as follows

Out[17]:

	feature	bow
0	someth	1
1	teach	1
2	later	1
3	thirti	1
4	day	1
5	rememb	1
6	tradi	1
7	air	1
8	show	1
9	book	1
10	see	1
11	ago	1
12	preschool	1
13	whole	1
14	use	1
15	televis	1
16	student	1
17	turn	1
18	seri	1
19	school	1
20	song	1
21	purchas	1
22	children	1
23	child	1
24	bought	1


Training the NB model using the train set and testing over the CV set with default values

```
In [19]: pred,acc=train(x_tr,y_tr,x_cv,y_cv)
```

The train accuracy by using default alpha over cv set is = 81.701449%

```
In [26]: #Confusion metric of the above model:
Confusion_metric(y_cv,pred,acc)
```

```
[[ 8563  4486]
 [ 9508 53919]]
```



+-----+   The performance metrics of the above model are as follows:   +-----+	
Metrics	Scores
Classification_accuracy	81.70144882054501
Classification_error	18.298551179454993
True positive	53919
False positive	4486
True negative	8563
False negative	9508
True positive rate	85.00953852460309
False negative rate	14.990461475396913
True negative rate	65.62188673461567
False positive rate	34.37811326538432
Precision value	92.31915075764061
Recall value	85.00953852460309
f1_score value	88.51369098430625
+-----+	

## OBSERVATIONS

- The train accuracy using default parameters of the model is 81.70% which is not that good for a classification model.
- Here the TP is dominating metric as compared to other metric which also results in high TPR value.
- Due to the low TN value the TNR value is very less (65.52%) which is not a good sign for a classification model.
- The precision, recall and the f1\_score cannot be trusted due to high TPR & low TNR values.


## Hyperparameter tuning the alpha value by using 10-fold Cross-validation technique

In [27]: `Score,mutinom =cross_validation(x_tr,y_tr)`

[0.7844576677081012, 0.7847770971595364, 0.7819526695323858, 0.7803667330368377, 0.781028009694429, 0.780030493409807, 0.7807029745481355, 0.7839196792801396, 0.7794869039976551, 0.7788928809182714]

In [28]: `optimal_a=MSE_plot(Score,mutinom) #CODE FOR PLOTTING THE ERROR PLOT`

The optimal number of alpha value is 0.000100%.



the misclassification error for alpha value is : [0.216 0.215 0.218 0.22 0.219 0.22 0.219 0.216 0.221 0.221]

## Testing the model by using the optimal alpha over the test set

In [29]: `#Testing the model with default value of alpha over the Test data  
test(x_tr,y_tr,x_test,y_test)  
  
print(" "*100)  
  
#Testing the model with optimal alpha value over the test data  
y_pred,Test_acc,=optimal_test(optimal_a,x_tr,y_tr,x_cv,y_cv,x_test,y_test)`

\*\*\*\* The Test accuracy by using default alpha is 81%  
\*\*\*\*\*

The train accuracy of the NB classifier for the best alpha=0.000100% is 81.633454%  
\*\*\*\*\*


The Test accuracy of the NB classifier by using the best alpha = 0.000100% is 81.199429%

## Confusion metric of the above model.



```
In [30]: Confusion_metric(y_test,y_pred,Test_acc)
```

```
[[12227 6854]
 [13686 76485]]
```



+-----+   The performance metrics of the above model are as follows:   +-----+		
	Metrics	Scores
+-----+		
	Classification_accuracy	81.1994288434079
	Classification_error	18.8005711565921
	True positive	76485
	False positive	6854
	True negative	12227
	False negative	13686
	True positive rate	84.82217120803807
	False negative rate	15.17782879196194
	True negative rate	64.07945076253864
	False positive rate	35.920549237461344
	Precision value	91.77575924837112
	Recall value	84.82217120803807
	f1_score value	88.16206558699788
+-----+		

OBSERVATION

- After training the model over optimal alpha value the model score decreased by 0.51% which is not good.
- This model scores are also similar to the previous model except there is slight decrease is seen on all the parameters.
- Since the model is trained over imbalanced dataset so the it is suffering from the high bias problem.
- So the performance can be increased if data balancing is performed over the dataset which can be done by using the oversampling technique.

Oversampling the data by using the SMOTE technique

```
In [31]: #FUNCTION FOR IMPLEMENTING THE SYNTHETIC MINORITY OVERSAMPLING TECHNIQUE
from imblearn.over_sampling import SMOTE

def Bal_train (X_tr, y_tr,X_cv,y_cv,X_test,y_test):

    sm = SMOTE()
    X_Train_res, y_Train_res = sm.fit_sample(X_tr, y_tr)
    X_Cval_res,y_Cval_res=sm.fit_sample(X_cv,y_cv)
    X_Test_res,y_Test_res=sm.fit_sample(X_test,y_test)

    clf = MultinomialNB()
    clf.fit(X_Train_res, y_Train_res )

    pred= clf.predict(X_Cval_res)

    acc = accuracy_score(y_Cval_res, pred, normalize=True) * float(100)

    print('\n The TRAIN accuracy by using default alpha over CV set is =  %f%% ' % ( acc))

    return X_Train_res, y_Train_res,X_Cval_res,y_Cval_res,X_Test_res,y_Test_res,pred, acc
```


Training the SMOTE model

```
In [34]: bal_xtr,bal_ytr,bal_xcv,bal_ycv,bal_xtes,bal_ytes,b_pre,b_acc=Bal_train(x_tr,y_tr,x_cv,y_cv,x_test,y_t
est)
```

The TRAIN accuracy by using default alpha over CV set is = 76.550996%

```
In [35]: Confusion_metric(bal_ycv,b_pre,b_acc)
```

```
[[45707 17720]
 [12026 51401]]
```



+-----+   The performance metrics of the above model are as follows:   +-----+	
Metrics	Scores
Classification_accuracy	76.55099563277469
Classification_error	23.449004367225314
True positive	51401
False positive	17720
True negative	45707
False negative	12026
True positive rate	81.03962035095465
False negative rate	18.960379649045358
True negative rate	72.06237091459474
False positive rate	27.93762908540527
Precision value	74.36379682006915
Recall value	81.03962035095465
f1_score value	77.5583184959411
+-----+	


## HYPERAMETER TUNING THE ALPHA FOR FINDING THE OPTIMAL ALPHA VALUE

```
In [36]: BScore,Mutinom =cross_validation(bal_xtr,bal_ytr)

[0.83896452215018, 0.8389808128280135, 0.8388504880422468, 0.8387885845704425, 0.8388374561793419, 0.8387820682143887, 0.8388179073234815, 0.8389319412191141, 0.8387592616051028, 0.8387462288929954]
```

```
In [37]: Boptimal_a=MSE_plot(BScore,Mutinom) #CODE FOR PLOTTING THE ERROR PLOT
```

The optimal number of alpha value is 0.000100%.



the misclassification error for alpha value is : [0.161 0.161 0.161 0.161 0.161 0.161 0.161 0.161 0.161 0.161]

### Testing the model by using the optimal alpha over the test set

```
In [38]: #Testing the model with default value of alpha over the Test data
test(bal_xtr,bal_ytr,bal_xtes,bal_ytes)

print("*"*100)

#Testing the model with optimal alpha value over the test data
Y_pred,BalTest_acc,=optimal_test(Boptimal_a,bal_xtr,bal_ytr,bal_xcv,bal_ytv,bal_xtes,bal_ytes)
```


\*\*\*\* The Test accuracy by using default alpha is 76%  
\*\*\*\*\*

The train accuracy of the NB classifier for the best alpha=0.000100% is 76.275876%  
\*\*\*\*\*

The Test accuracy of the NB classifier by using the best alpha = 0.000100% is 75.777689%

```
In [39]: Confusion_metric(bal_ytes,Y_pred,BalTest_acc)
```

```
[[63453 26718]
 [16965 73206]]
```



+-----+-----+		
The performance metrics of the above model are as follows:		
+-----+-----+		
	Metrics	Scores
+-----+-----+		
	Classification_accuracy	75.7776890574575
	Classification_error	24.222310942542503
	True positive	73206
	False positive	26718
	True negative	63453
	False negative	16965
	True positive rate	81.18574708054696
	False negative rate	18.81425291945304
	True negative rate	70.36963103436803
	False positive rate	29.630368965631966
	Precision value	73.26167887594572
	Recall value	81.18574708054696
	f1_score value	77.02043714984613
+-----+-----+		

## OBSERVATION

- After oversampling the data the accuracy of the model reduced (from 81% to 75%) as compared to the previous model but the model is very stable and sensible because of high TP and TN values.
- There is a considerable increase in the TNR value (64%-70.36%)& considerable decrease in the FPR value(34% to 29.63%) which is a good sign for the improvement of the model.
- So the model is sensible and stable as compared to the previous model which is achieved by balancing the data.
- The metric as accuracy is very misleading and can't be trusted in a imbalanced data.

## Implemnting the Tf-idf vectorization technique

```
In [67]: #Initializing the count vectorizer
TF_vect=TfidfVectorizer(ngram_range=(1,2),binary=True)

#vectorizing the X_train set
TF_count,X_tr=vec_train(TF_vect,X_tr["CleanedText"])

print("The shape of the X_train is: ",X_tr.shape)

#Vectgorizing the X_crossvalidation set
X_cv=vec_cv(TF_count,X_cv["CleanedText"])
print("The shape of the X_cv is: ",X_cv.shape)

#Vectorizing the X_test set
X_test=vec_test(TF_count,X_test["CleanedText"])
print("The shape of the X_test is: ",X_test.shape)

#Printing the total length of the features
print("\nTop 25 feaures according to the Bow score are as follows")
Features = TF_vect.get_feature_names()
len(Features)

top_tfidf = top_tfidf_feats("tfidf",X_tr[1,:].toarray()[0],Features,25)
top_tfidf
```

The shape of the X\_train is: (178443, 1816552)

The shape of the X\_cv is: (76476, 1816552)

The shape of the X\_test is: (109252, 1816552)

Top 25 feaures according to the Bow score are as follows

Out[67]:

	feature	tfidf
0	preschool turn	422.426325
1	teach preschool	422.426325
2	whole school	422.426325
3	school purchas	422.426325
4	sister later	422.426325
5	song student	422.426325
6	air televis	422.426325
7	child sister	422.426325
8	book song	422.426325
9	use seri	372.605606
10	seri book	366.076962
11	book children	358.410533
12	televis year	326.201503
13	day thirti	321.245669
14	thirti someth	310.170360
15	children tradit	308.586244
16	show air	297.616678
17	student teach	282.566459
18	ago child	246.284028
19	along book	203.244144
20	see show	196.984144
21	later bought	161.464536
22	tradi live	145.680425
23	turn whole	137.896025
24	rememb see	112.189152

Training the Tf-idf vectorized model over the CV set.


```
In [41]: Predi,Acc=train(X_tr,y_tr,X_cv,y_cv)

The train accuracy by using default alpha over cv set is = 85.608557%
```

Confusion metric of the above model.

```
In [42]: Confusion_metric(y_cv,Predi,Acc)

[[ 4064  8985]
 [ 2021 61406]]
```



+-----+   The performance metrics of the above model are as follows:   +-----+	
Metrics	Scores
Classification_accuracy	85.60855693289398
Classification_error	14.39144306710602
True positive	61406
False positive	8985
True negative	4064
False negative	2021
True positive rate	96.81365979787788
False negative rate	3.1863402021221248
True negative rate	31.1441489769331
False positive rate	68.8558510230669
Precision value	87.23558409455754
Recall value	96.81365979787788
f1_score value	91.77539643396254
+-----+	

## OBSERVATION

- The accuracy of the model by using the default alpha value is 85.60% which is good but the other metrics are very less and very alarming.
- In this model the TP value is very dominating which results in high TPR rate as compared to the other metrics which is not good as the model will become biased towards positive reviews.
- The TNR is very less and this model will fail to classify the negative reviews properly.
- Let's see if the metrics become better after doing hyper-parameter tuning the alpha value.


## HYPERAMETER TUNING THE ALPHA FOR FINDING THE OPTIMAL ALPHA VALUE

```
In [43]: CV_Score,Mutinomi =cross_validation(X_tr,y_tr)

[0.8456538275213425, 0.850350012363933, 0.7999472979891289, 0.7642888248070557, 0.781392345001281, 0.7573846559978202, 0.7729134422929673, 0.8363399278287986, 0.742136071851881, 0.7267473805074696]
```

```
In [44]: Optimal_a=MSE_plot(CV_Score,Mutinomi) #CODE FOR PLOTTING THE ERROR PLOT
```

The optimal number of alpha value is 0.000100%.



the misclassification error for alpha value is : [0.154 0.15 0.2 0.236 0.219 0.243 0.227 0.164 0.258 0.273]

## Testing the model by using the optimal alpha over the test set

```
In [45]: #Testing the model with default value of alpha over the Test data
test(X_tr,y_tr,X_test,y_test)

print("***100)

#Testing the model with optimal alpha value over the test data
Y_prediction,Test_accuracy=optimal_test(Optimal_a,X_tr,y_tr,X_cv,y_cv,X_test,y_test)

**** The Test accuracy by using default alpha is 84%
*****


The train accuracy of the NB classifier for the best alpha=0.000100% is 84.246038%
*****

The Test accuracy of the NB classifier by using the best alpha = 0.000100% is 83.790686%
```

## Confusion metric of the above model.

```
In [46]: Confusion_metric(y_test,Y_prediction,Test_accuracy)

[[ 3501 15580]
 [ 2129 88042]]
```



+-----+-----+		
The performance metrics of the above model are as follows:		
+-----+-----+		
	Metrics	Scores
+-----+-----+		
	Classification_accuracy	83.79068575403655
	Classification_error	16.209314245963462
	True positive	88042
	False positive	15580
	True negative	3501
	False negative	2129
	True positive rate	97.63893047653902
	False negative rate	2.36106952346098
	True negative rate	18.348094963576333
	False positive rate	81.65190503642367
	Precision value	84.96458281060006
	Recall value	97.63893047653902
	f1_score value	90.8618990366009
+-----+-----+		

## Observations

- The accuracy of the model has become low after tuning the alpha which is not good for a classification model.
- The main reason is because of the high FPR and low TNR which affects the model a lot.
- This model is also suffering from the heavy bias problem and may be solved by oversampling the datapoints.
- The high Precision,recall and f1\_score value cannot be trusted because of high tpr as compared to other metrics.

## Oversampling the data by using the SMOTE technique


```
In [47]: Bal_xtr,Bal_ytr,Bal_xcv,Bal_ycv,Bal_xtes,Bal_ytes,B_pre,B_acc=Bal_train(X_tr,y_tr,X_cv,y_cv,X_test,y_t
est)
```

The TRAIN accuracy by using default alpha over CV set is = 62.409542%

## Confusion metric of the above model.

```
In [48]: Confusion_metric(Bal_ycv,B_pre,B_acc)
```

```
[[17149 46278]
 [ 1407 62020]]
```





+-----+-----+		
The performance metrics of the above model are as follows:		
+-----+-----+		
	Metrics	Scores
+-----+-----+		
	Classification_accuracy	62.409541677834355
	Classification_error	37.59045832216564
	True positive	62020
	False positive	46278
	True negative	17149
	False negative	1407
	True positive rate	97.78170179891845
	False negative rate	2.2182982010815584
	True negative rate	27.03738155675028
	False positive rate	72.96261844324971
	Precision value	57.26790891798556
	Recall value	97.78170179891845
	f1_score value	72.23176590478964
+-----+-----+		


## HYPERAMETER TUNING THE ALPHA FOR FINDING THE OPTIMAL ALPHA VALUE

```
In [49]: BScoreS,Mutinomial =cross_validation(Bal_xtr,Bal_ytr)

[0.9823736553630698, 0.9827581127274145, 0.9804546256578697, 0.9791383491218069, 0.9797704214348719,
0.978907023152515, 0.9794218046657244, 0.9818979690139731, 0.9783401129138823, 0.977802524154485]
```

```
In [50]: Bal_optimal_a=MSE_plot(BScoreS,Mutinomial) #CODE FOR PLOTTING THE ERROR PLOT
```

The optimal number of alpha value is 0.000100%.



the misclassification error for alpha value is : [0.018 0.017 0.02 0.021 0.02 0.021 0.021 0.018 0.022 0.022]

## Testing the model by using the optimal alpha over the test set

```
In [51]: #Testing the model with default value of alpha over the Test data
test(Bal_xtr,Bal_ytr,Bal_xtes,Bal_ytes)

print("***100)

#Testing the model with optimal alpha value over the test data
Y_PRED,BalTest_Acc=optimal_test(Bal_optimal_a,Bal_xtr,Bal_ytr,Bal_xcv,Bal_ycv,Bal_xtes,Bal_ytes)
```

```
**** The Test accuracy by using default alpha is 61%
*****
```


```
The train accuracy of the NB classifier for the best alpha=0.000100% is 57.949296%
*****
```

The Test accuracy of the NB classifier by using the best alpha = 0.000100% is 57.969857%

## Confusion metric of the above model.

```
In [52]: Confusion_metric(Bal_ytes,Y_PRED,BalTest_Acc)
```

```
[[16221 73950]
 [ 1848 88323]]
```



+-----+   The performance metrics of the above model are as follows:   +-----+	
Metrics	Scores
Classification_accuracy	57.96985727118474
Classification_error	42.03014272881525
True positive	88323
False positive	73950
True negative	16221
False negative	1848
True positive rate	97.95056060152378
False negative rate	2.0494393984762285
True negative rate	17.989153940845725
False positive rate	82.01084605915428
Precision value	54.42864801907896
Recall value	97.95056060152378
f1_score value	69.97433094072348
+-----+	

## Observations:

- Since after doing the data balancing technique the model score decreased very drastically due to poor performance metrics.
- I think the model is overfitting which leads to decrease in accuracy and other factor may be the high dimensions of the tf-idf vectorized data.
- In this model SMOTE didn't improved the model in terms of stability and accuracy.

In [ ]: